

AI News Aggregator — Product & Functionality Spec

2026-09-18 · @Someone

What This Is

An AI-focused news site that aggregates, deduplicates, and re-cuts coverage from both AI-specialist newsletters and mainstream press into a single browsable, filterable feed. Rather than reporting news first-hand, it acts as a second-order aggregator: pulling from sources that already do the sourcing work (Rundown AI, TLDR AI, Reuters, FT, SCMP, etc.) and adding a layer of classification and editorial judgment that no existing aggregator offers.

Who It's For

- General AI-curious readers who want broad coverage without subscribing to a dozen newsletters
- HR, L&D, org design, and change management professionals tracking AI's effect on work (a plausible paid-tier audience)
- Readers specifically interested in AI's effect on daily life, ethics, and global equity — audiences currently underserved by existing aggregators

What Makes It Different

Existing aggregators (Feedly, Digg, Particle, Readless, The Rundown AI, TLDR AI) are strong on speed and volume but tone-neutral, topically generic, and geographically US-centric. None of them let a reader filter by:

- **Tone** (Good / Useful / Bad / Ugly) — an explicit point of view on whether a story represents progress, practical value, failure, or harm
- **Topic**, including verticals nobody else covers as first-class categories: Future of Work and Future of Daily Life
- **Geography**, beyond the default US-skew of English-language coverage — including Global South / underrepresented-region coverage
- **Access** — whether the underlying story is free to read or sits behind a paywall

The combination of these four axes, applied consistently across every story, is the product's defensible niche.

Format

A full self-serve website with browse-and-filter functionality (not a newsletter-only product), so that once live, reading requires no ongoing manual curation from the operator — supporting a largely passive-income model. A digest/newsletter version can be generated from the same tagged data as a secondary, low-effort product surface.

Curation Model

An AI pipeline ingests, deduplicates, and classifies stories automatically. The operator does not review every story — only a spot-check queue of low-confidence classifications (see Pipeline Functionality), keeping the ongoing time investment low while preserving editorial quality on the judgment-heavy calls (e.g. Future of Work vs. generic Tools & Products; Ethics vs. Policy & Regulation).

Content Taxonomy

Every story is tagged with one or more topics. Topics are what the story is *about* — independent of tone, access, or geography (see next section).

Topic	Covers	Distinguished from
Models & Research	New model releases, benchmarks, research breakthroughs	Tools & Products (this is the underlying tech, not the shipped tool)
Business & Funding	Funding rounds, M&A, market moves, company earnings	Policy & Regulation (this is market activity, not government action)
Policy & Regulation	Laws, enforcement actions, court cases, government rules constraining or governing AI	National Initiatives (this restricts AI; National Initiatives builds AI capacity)
National Initiatives	Sovereign AI strategy, government investment in compute/data centers, state-backed labs, national talent/skills programs, industrial policy	Policy & Regulation (this is investment/strategy, not legal constraint)
Ethics	Bias, safety incidents, IP/copyright disputes, misuse, surveillance, deepfakes, transparency failures	Policy & Regulation (this is contested regardless of whether law has caught up)

Topic	Covers	Distinguished from
Future of Work	Upskilling, training, process re-engineering, change management, org design, job displacement/creation, workplace AI adoption	Future of Daily Life (this is organizational/professional, not personal/consumer)
Future of Daily Life	How AI shows up outside work: relationships, education, healthcare, creativity, dating, social media, mental health, everyday consumer tools. Includes a Children & Education subtopic (AI in classrooms, child safety, parenting)	Future of Work (this is personal/consumer, not organizational)
AI Equity & Representation	Underrepresentation in training data (languages, cultures, dialects), the digital divide, talent drain, local/sovereign AI capacity-building in underserved regions, who AI is built for and who it excludes	National Initiatives (this is about exclusion/inclusion; National Initiatives is about a given nation's own strategy)
Tools & Products	Launches, how-to guidance, practical use of shipped AI products	Models & Research (this is the shipped product, not the underlying research)

A story can carry more than one topic tag (e.g. a story about an EU-funded reskilling program is both National Initiatives and Future of Work).

Classification Axes

Every story carries a topic tag (above) plus three independent axes. A story can be, for example: Topic = Ethics, Tone = Ugly, Access = Free, Geography = US — all four assigned separately.

Tone

Tone	Definition	Example
Good	Genuine progress or positive outcome	A new model materially improves medical diagnosis accuracy
Useful	Practical, actionable value for the reader	A how-to on using an AI tool for a specific task
Bad	Failure, backlash, hype-vs-reality gap	A high-profile AI product underdelivers on its claims
Ugly	Harm, safety incident, scandal	A deepfake scandal, a safety researcher's resignation over a deployment

Access

Access	Definition
Free	Story is fully readable without a subscription or paywall
Paid	Story sits behind a paywall (e.g. FT, Bloomberg, WSJ) — flagged so readers know before they click

Geography

Geography	Covers
Worldwide	No single-country focus, or a global trend/comparison
US	United States-specific developments
China	China-specific developments
Europe	EU and wider European developments
Africa	Africa-specific developments
Latin America	Latin America-specific developments
South & Southeast Asia	India, Southeast Asia-specific developments
Middle East	Gulf states and wider Middle East developments

The last four regions start with lower expected volume than US/China/Europe and are prioritized specifically because the AI Equity & Representation topic depends on them existing as real filters, not a generic "Rest of World" catch-all. Volume in each will be monitored and the list adjusted post-launch.

Content Type

A fifth axis, alongside Topic, Tone, Access, and Geography — what kind of content this is, independent of what it's about:

Content Type	Covers
Article	A written news story (the default type everything above was written around)
Podcast	An audio episode discussing AI — classified from its show notes/description, not a full transcript (see Pipeline Functionality)
Video	Any video content about AI — frequent YouTube explainer/news channels (ingested automatically like podcasts) and rare one-off documentaries (curated manually) share this type; see Source Map for how each is actually sourced

Topic, Tone, Access, and Geography still apply to podcast and documentary items exactly as they do to articles — a podcast episode can be tagged Future of Work + Useful + Free + Worldwide the same way an article would be.

Source Map

Sources split into two kinds: **AI-specialist sources** (already do the sourcing/curation work) and **mainstream/general news outlets** (broader reach, but AI is one topic among many — needed so the aggregator doesn't miss stories that mainstream press covers before AI-specialist newsletters do, and so mainstream framing of AI stories — e.g. labor, education, politics desks — is represented alongside tech-press framing).

AI-Specialist Sources (feed most topics)

Source	Feeds	Access
The Rundown AI	General/Worldwide	Free

Source	Feeds	Access
TLDR AI	General/Worldwide	Free
The Neuron	General/Worldwide	Free
Techpresso	General/Worldwide	Free
The Batch (Andrew Ng)	Models & Research	Free
AI Weekly	General, significance-ranked, includes a China tracker	Free
Turing Post	General, with dedicated coverage of China/Israel/Europe	Free
Digg (AI edition)	General, AI + human editorial	Free

Future of Work Sources

Source	Notes
One Useful Thing	Closest existing analog to this vertical; how AI changes work, school, life
The Median (DataCamp)	Connects AI/data developments to skills and upskilling
Chain of Thought	AI and creativity, agentic AI reshaping workforce norms
SHRM, Gallup, WEF Future of Jobs Report, LinkedIn Work Change Report	Institutional research (not newsletters) — rich source material, lower frequency
McKinsey ("Lead with AI"), BCG, Gloat, Dayforce	HR-tech and consulting research on workforce/upskilling trends

China AI Sources

Source	Notes	Access
AI Weekly's China AI tracker	Aggregates FT, Bloomberg, Reuters, TechCrunch, Wired, Nikkei coverage of Chinese labs/models	Free (tracker); underlying sources often Paid

Source	Notes	Access
South China Morning Post	Direct regional coverage	Free (metered)
Bloomberg, FT, Reuters (China desks)	Primary reporting	Paid

Editorial standard for China stories: verify claims against the original technical report, GitHub repository, or official lab post before publishing — English-language summaries of Chinese AI releases are frequently distorted in translation or inflated for investor audiences.

Europe / EU AI Policy Sources

Source	Notes	Access
EU AI Act Explorer + newsletter (Future of Life Institute)	Authoritative free tracker, 50,000+ subscribers, bi-weekly	Free
European Commission — Shaping Europe’s Digital Future	Primary source press releases	Free
AI in Europe	"Full picture of European AI, beyond Brussels" — enforcement, funding, policy combined	Free
Law firm trackers (White & Case, DLA Piper, Littler)	Useful for the Ethics/regulatory-risk angle	Free

AI Equity & Representation / National Initiatives Sources

Weaker existing coverage here than other topics — likely requires direct monitoring of:

- National government AI strategy pages (e.g. India’s IndiaAI Mission, UK compute strategy, Gulf state sovereign AI investment announcements)
- Regional tech press (e.g. Rest of World, TechCabal for Africa, Rest of World’s Latin America and Southeast Asia coverage)
- Multilateral/NGO reports (UN, World Bank, ITU digital divide research)

Mainstream / General News Sources (non-AI-focused, but covering AI)

Source	Why included	Access
Reuters, AP, Bloomberg, FT	Break major AI business/policy stories often before AI-specialist newsletters	Mixed (FT/Bloomberg largely Paid)
The Guardian, NYT, WaPo	Strong "AI and society/daily life" framing, useful for Future of Daily Life and Ethics	Mixed (metered/Paid)
The Economist	Strong policy and economics framing, useful for National Initiatives and Business & Funding	Paid
Local/regional outlets (SCMP for China, Africa/LatAm regional press)	Needed for non-US-centric coverage and the Equity topic specifically	Mixed

Coverage Gaps to Note

AI Equity & Representation and National Initiatives (outside the US/China/EU) currently have the thinnest source base of any topic in the taxonomy. These sections may update weekly rather than daily at launch until dedicated regional sources are identified.

Podcast Sources

Source	Notes
Dwarkesh Podcast	Long-form frontier-lab interviews
Latent Space	AI engineering-focused, weekly
The Cognitive Revolution	Strategy-oriented, long-form
No Priors	Founder/investor perspective
The AI Daily Brief	Daily news-briefing format — closest analog to the article sources above
Machine Learning Street Talk	Technical depth
The TWIML AI Podcast	Broad ML/AI research and industry coverage
NVIDIA's The AI Podcast	Applied AI across industries, biweekly

All have public RSS feeds with episode titles, show notes, and descriptions — sufficient for classification without needing full audio transcription. Transcription (e.g. via a speech-to-text API) would allow deeper summarization but is a meaningfully bigger cost and complexity addition, and is treated as a future enhancement rather than an MVP requirement.

Note on Spotify specifically: most shows above are also distributed on Spotify, but that's irrelevant to ingestion — what matters is whether a show has an open RSS feed, which most independent podcasts do regardless of where they're also listened to. Spotify-exclusive shows with no public RSS are a genuine gap: reaching those would need Spotify's own API, which is built for personalized playback rather than full-catalog indexing and would likely require partner-level access rather than a simple feed pull. Not worth building around at MVP — flag it as a known gap rather than a blocker.

YouTube AI Channels

Source	Notes
Two Minute Papers	Short, frequent research-paper breakdowns
AI Explained	Frontier AI developments and analysis
Matt Wolfe	Weekly AI news and tool coverage
Wes Roth	AI news and breakthroughs

Source	Notes
Fireship	Fast-paced tech/AI explainers
Yannic Kilcher	ML paper reviews

Every YouTube channel exposes a free public RSS feed of its recent uploads (no API key required), so these are ingested on the same schedule and classified from title/description the same way as the podcast sources above — genuinely different from the one-off documentaries below, which have no feed to poll.

One-off Documentaries

AI-focused documentaries and investigative long-form video are released rarely and irregularly — there’s no equivalent to an RSS feed producing a steady stream of these. This category is realistically a **manually curated addition**: the operator adds a documentary or long-form piece to the site directly (through the same interface as the spot-check queue) when one is released, rather than the automated pipeline discovering it. Likely sources to monitor periodically: major streaming documentary units (Netflix, Bloomberg Originals), broadcast investigative strands (PBS Frontline, BBC), and long-form YouTube channels doing AI explainers or investigations. Given the low volume, this content type should be expected to update monthly or less, not daily.

Story Deduplication & Grouping

The same underlying story is routinely covered by 5–30 different sources within hours (a model launch, a policy announcement). Without deduplication, the feed becomes unusable noise — this is a baseline requirement, not an enhancement. This deduplication logic applies to Article-type content, where the same underlying fact genuinely is reported by many outlets. Podcast episodes and documentaries are original content in their own right, not competing reports of the same event, so they are listed individually rather than merged — even when two podcasts happen to discuss the same news story in the same week.

Detection Approach

- 1. **Candidate matching** — for each incoming item, generate an embedding of its headline + first paragraph and compare against items ingested in the last 48–72 hours. Items above a similarity threshold become dedup candidates (cheap, fast, catches most duplicates).

2. **Entity/fact overlap check** — cross-check candidates on shared named entities (company, model name, person, dollar figure, date) extracted from each item. This catches near-duplicates that phrase things differently but report the same underlying fact, and filters out false positives from embedding similarity alone (two different funding rounds for competing companies can read as "similar" on embeddings but shouldn't merge).
3. **LLM confirmation for borderline cases** — items that pass step 1 but are ambiguous on step 2 get a direct LLM check: "are these the same underlying event?" Only ambiguous cases hit this step, to control cost.

Grouping Behavior

- Confirmed duplicates merge into **one story card** on the site, not one per source.
- The merged card carries: one summary (written fresh, not copied from any single source — see Copyright Note below), a list of every contributing source with its own link, and the tags (topic/tone/access/geography) that most sources agree on.
- If sources disagree on tags (e.g. one outlet frames a story as "Good," another as "Bad"), the story is tagged with the majority view and flagged for a spot-check rather than auto-resolved silently — tone disagreement is exactly the kind of judgment call that shouldn't be fully automated.
- If sources disagree on access (some free, some paid), the card shows both, so a reader always has at least one way to read further even if their preferred source is paywalled.
- A later development of an existing story (a follow-up, a correction, new information) is linked to the original story as an update, not merged into it as a duplicate and not treated as an unrelated new story.

Copyright Note

Summaries are written fresh from the underlying facts, never copied or lightly reworded from any single source, and no single source is quoted at length — this keeps the aggregator legally sound and avoids the display-page-than-substance problem that weaker aggregators have.

Pipeline Functionality

```

flowchart LR
  A[Ingest] --> B[Dedupe & Group]
  B --> C[Classify & Tag]
  C --> D{Confidence check}
  D -->|High confidence| E[Publish]
  D -->|Low confidence| F[Spot-check queue]
  F --> E
  E --> G[Digest generation]

```

1. **Ingest** — pull from AI-specialist sources, mainstream outlets, and institutional research feeds (RSS where available, scraping where not) on a rolling schedule (e.g. every 30–60 minutes for daily-news sources, daily for slower institutional sources). Before scraping any source without an RSS feed or official API, the worker checks that source's robots.txt and stated terms of service for a text-data-mining opt-out; sources that opt out are skipped rather than scraped around, per the Legal & Compliance Considerations below.
2. **Dedupe & Group** — as described above; produces one grouped story per underlying event.
3. **Classify & Tag** — an LLM call per grouped story assigns topic(s), tone, access, geography, and writes the fresh summary.
4. **Confidence check** — the classifier also outputs a confidence score per tag. Low-confidence calls are routed to a spot-check queue rather than published directly; high-confidence calls publish automatically. The fuzziest boundaries in this taxonomy — Future of Work vs. Tools & Products, and Ethics vs. Policy & Regulation — are expected to generate the most spot-check volume early on, and the classifier prompts should be refined as patterns emerge. Podcast episodes are ingested from their RSS feeds on the same schedule as article sources, classified from show notes/descriptions rather than audio. YouTube AI channels are ingested the same way, via each channel's public RSS feed. One-off documentaries are added directly by the operator (see Source Map) rather than through this automated step, since there's no reliable feed to poll for those. A generated summary that scores too similar to its source text (a basic text-similarity check against the original) is also routed to spot-check rather than published automatically — a functional guardrail against the fresh-summary requirement drifting into a lightly-reworded copy, which matters given the EU press publishers' right discussed below.
5. **Spot-check queue** — the operator reviews only flagged items, not the full feed. A review interface shows the story, its proposed tags, and a one-click approve/correct action.
6. **Publish** — approved and high-confidence stories go live on the site.

7. **Digest generation** — a daily and/or weekly email digest is auto-compiled from the same published, tagged data (e.g. "today's Future of Work stories" or "this week's Ethics + Ugly"), giving a secondary product surface at near-zero incremental effort.

Site Functionality


- **Home feed** — reverse-chronological grouped stories, each shown once regardless of how many sources covered it.
- **Filters** (combinable) — Topic, Tone, Access, Geography. Content Type (Article / Podcast / Video) is a fifth filter, combinable with the rest. A reader can select any combination (e.g. Future of Work + Ugly + Europe).
- **Story card** — fresh summary, topic/tone/geography/access tags, list of contributing sources with individual links, and (if applicable) a link to a later update on the same story. Carries a visible "AI-generated summary" note (an AI Act Article 50 transparency requirement, not just good practice) and a "report an issue" link for readers to flag a factual error or mistagging, which routes to the spot-check queue rather than a public comment thread. A podcast or video card additionally shows a thumbnail, running time, and an embedded player or a direct link to listen/watch, in place of the source list an article card shows.
- **Search** — free-text search across published summaries.
- **Saved filter views** — a reader can save a filter combination (e.g. "Future of Work, Worldwide") as a personal default view or a subscribed digest.
- **Digest sign-up** — email capture tied to a filter combination, feeding the digest generation step in the pipeline above. Sign-up uses an explicit, non-pre-ticked consent checkbox and links the privacy policy, per GDPR.
- **Paid tier (future)** — candidate gated features: deeper Future of Work coverage, an API/export of tagged data, or ad-free browsing. Not required for MVP launch. Whenever this is built, checkout needs to capture two non-contradictory pieces of customer location evidence (billing address, card country, or IP), since that's what EU VAT rules require to charge the right country's rate — see VAT (EU One-Stop-Shop) below.

Open Questions & Next Steps

- ☐ Decide MVP scope: full taxonomy at launch, or a reduced set (e.g. drop AI Equity & Representation and National Initiatives until source volume is proven) with expansion in v2
- ☐ Identify and confirm concrete source feeds (RSS availability, scraping legality/ToS) for each topic, especially the thin Equity/National Initiatives category

- ☐ Define the specific confidence threshold and spot-check cadence (daily review session length)
- ☐ Decide tech stack and hosting for the ingestion pipeline and the site
- ☐ Decide free vs. paid tier structure and pricing
- ☐ Decide digest cadence and format (daily vs. weekly, per-topic vs. general)

Architecture Proposal



Recommended Stack

Bias toward managed/serverless services throughout — the goal is to minimize ongoing DevOps burden given this is a side project, not a team with infrastructure headcount.

Layer	Recommendation	Why
Frontend (site UI)	Lovable, generating the filter UI, story cards, search, and spot-check review page	Fastest path to a working site without hand-coding, and pairs natively with Supabase as its backend — no separate stack to wire together
Database	Supabase (managed Postgres)	Lovable's default backend; the same database the ingestion/classification worker (below) writes into, so frontend and pipeline share one source of truth without any custom syncing

Layer	Recommendation	Why
Ingestion, dedup & classification	A separate small always-on worker (e.g. a €5–9/month host like Fly.io, Railway, or Render) running scraping, dedup, and the Haiku classification calls on a schedule, writing results to the shared Supabase database	Kept outside Lovable deliberately — scraping messy mainstream sites and running embedding-based dedup across hundreds of daily items is exactly the kind of custom backend logic that doesn't map cleanly onto Lovable's frontend-first model or fit inside a lightweight scheduled function
Deduplication	Embeddings (a small/cheap embedding model) + entity extraction, run inside the ingestion worker	Keeps the expensive step (LLM classification) reserved for content that's passed dedup, not run on every raw item
Classification & tagging	Claude Haiku via the Anthropic API, using the Batch API for the bulk of calls, run inside the ingestion worker	Haiku is the cheapest current Claude tier suited to classification/routing/short-summary work; batch processing cuts token cost by roughly half
Spot-check queue	A review page built as part of the Lovable frontend, reading/writing the same Supabase tables	No separate admin tool needed — it's just another view Lovable generates against the shared database
Digest/email	Resend or Beehiiv	Both have workable free tiers for a small subscriber base and handle deliverability, which is not worth building in-house

Why This Shape

The expensive part of any AI pipeline is unconstrained LLM usage — calling a large model on every raw item, with no caching or batching. This architecture avoids that by design: cheap embedding-based matching filters out most duplicates before anything reaches the LLM, and the LLM step itself runs on the cheapest capable model, in batch mode, once per grouped story rather than once per source. That ordering is the main lever on the cost figures in the next two sections.

Build Tool vs. Runtime Model

Worth keeping distinct: the tool used to **write** the pipeline's code is separate from the model that code **calls once running**. These don't need to match, and mixing them is normal:

- **Codex** (included in an existing ChatGPT Plus plan) is used interactively during the build — writing and debugging the scraper, dedup logic, and the code that calls the classification API. This costs nothing beyond the existing Plus subscription.
- **Claude Haiku via the Anthropic API** is what that code calls at runtime for the classify/tag/summarize step, billed separately per token (the ~€13-17/month estimate in Cost of Running, below). A ChatGPT Plus subscription does not cover this — API usage is billed independently of any consumer chat subscription, on either platform.

The practical implication for whoever builds this: the code Codex helps write should call the Anthropic API for classification, not an OpenAI model, even though the coding assistant itself is OpenAI's.

Cost of Building (One-Time)

Two realistic paths, with a wide range because scope (nine topics, four axes, dedup logic, spot-check UI) is genuinely substantial for a first version: (figures below converted at roughly \$1 ≈ €0.87, the September 2026 rate — treat as directional given currency movement)

Path	Estimated cost	Notes
DIY, using an AI coding assistant (e.g. Codex, already included in your ChatGPT Plus plan) to write the code	Mostly your own time; incidental API/tooling costs likely €45-175 total during the build	Realistic given your BA/PO background and the fact you've already built similar tools (the AI Transformation business-case tool). Time cost is the real cost here, not cash — plan for several weeks of part-time effort given the taxonomy's complexity

Path	Estimated cost	Notes
Freelance developer / small agency for an MVP	Roughly €3,500–€10,500, depending on how much of the filter UI, dedup pipeline, and spot-check tooling is custom-built vs. assembled from existing services	Wide range reflects genuine uncertainty — get quotes against a written spec (this document) before committing, and treat any single number a contractor gives you as provisional until they've seen the full taxonomy

A middle path — DIY for the core pipeline and site, with a freelancer brought in only for a specific hard part (e.g. the dedup logic or the classification prompts) — is worth considering if you hit a wall on one component rather than assuming an all-or-nothing build.

Either way, the full nine-topic taxonomy at launch means the build should budget real time for classifier prompt-writing and testing across all four axes, not just the pipeline plumbing — getting Future of Work vs. Tools & Products and Ethics vs. Policy & Regulation right (the two fuzziest boundaries flagged earlier) will likely take several rounds of prompt iteration against real story examples before the spot-check volume settles down.

Cost of Running (Monthly, at MVP Scale)

Assumptions: ~300–500 raw items ingested per day across all sources and topics (the full nine-topic taxonomy pulls from more sources than a narrower version would), deduplicating down to roughly 100–150 grouped stories per day published to the site.

LLM Classification Cost (the AI pipeline itself)

Using Claude Haiku 4.5 via the Batch API (current published rate: €0.90/€4.50 per million input/output tokens standard, roughly half that with batch processing — so approximately €0.45/€2.25 per million tokens batched):

Step	Per-story estimate	Daily (150 stories)	Monthly
Classify + tag + summarize (one call per grouped story, ~2,000–3,000 input tokens covering all contributing sources, ~300–400 output tokens for tags + fresh summary)	~€0.002–0.003	~€0.30	~€9

Step	Per-story estimate	Daily (150 stories)	Monthly
Embedding-based dedup matching (runs on raw items pre-grouping, cheap per-item)	~€0.0001–0.0003	~€0.09	~€3
Spot-check re-classification (occasional second pass on flagged items)	Marginal	~€0.04	~€1–2
AI pipeline subtotal			~€13–17/month

This is a genuinely small line item relative to hosting and services below — LLM classification at this scale and with Haiku-tier pricing is not the cost driver people often assume it will be. It would only grow meaningfully if story volume scales up an order of magnitude, or if a costlier model is used for every call rather than reserved for ambiguous cases.

Hosting & Services

Service	Estimated cost	Notes
Frontend (Lovable)	€0–22/month	Lovable bills on a credit/message system rather than pure traffic. The free tier covers light building; the Starter/Pro tier (~€17–22/mo) is likely once you're actively iterating or want a custom domain — a genuinely new line item the earlier Vercel-based plan didn't carry, since Vercel's free tier would have covered a site at this scale for €0
Database (Supabase)	€0–22/month	Free tier covers MVP scale; paid tier if story volume/history grows large
Ingestion worker hosting (Fly.io, Railway, or Render)	€5–9/month	Runs the scraping/dedup/classification code separately from Lovable, as described in Architecture Proposal

Service	Estimated cost	Notes
Email/digest (Resend or Beehiiv)	€0-17/month	Free tiers typically cover a few hundred to low-thousands of subscribers
Domain	€6.53/month (actual quote: €156.75 for 2 years, amortized)	
Anthropic API (see above)	~€13-17/month	
Total estimated monthly running cost	~€28-92/month , realistically closer to €42-55/month at launch — the domain is now a known, fixed €6.53/month rather than an estimate, and Lovable Starter + the ingestion worker remain the other two line items least likely to stay at €0	Wide range reflects free-tier headroom on most services; costs rise mainly with traffic/subscriber growth, not with the AI classification itself

These are current published rates and free-tier terms as of this writing; hosting and API pricing both change periodically, so treat this as directional rather than a locked-in budget.

Cost of Maintaining (Ongoing)

Maintenance here is mostly time, not money — the monthly cash figures above already assume the pipeline is running unattended most of the time.

Task	Estimated time	Frequency
Spot-check queue review	15-30 minutes	Most days, though this can slip to every 2-3 days without the site going stale, since only low-confidence items queue

Task	Estimated time	Frequency
Classifier prompt tuning	1-2 hours	Occasional, concentrated in the first 4-8 weeks post-launch while the Future of Work/Tools & Products and Ethics/Policy boundaries are being tuned; drops off sharply once patterns stabilize
Source health checks (a feed breaks, a scraper stops working)	30 minutes-2 hours	As needed; more likely for scraped mainstream sources than RSS feeds, and more likely for the thinner Equity/National Initiatives sources flagged earlier
Taxonomy or axis adjustments (e.g. splitting "Rest of World" into named regions as volume justifies it)	Variable	Quarterly review is a reasonable cadence

Given you want this to be genuinely passive, the practical target is: spot-check review as a short daily or every-other-day habit, everything else handled reactively rather than on a schedule. If the daily spot-check habit itself becomes the bottleneck, that's a signal to either tighten the confidence threshold (publish more automatically, accept slightly more tagging noise) or accept a lower daily story volume, rather than a sign the whole model has failed.

Breakeven Analysis

Subscriber count needed at a given monthly price to cover the running cost alone (not the one-time build cost, and not a return on your own time) — calculated against both the realistic launch estimate (€42-55/month) and the wider full-range estimate (€28-92/month) from Cost of Running, above. Updated to reflect the actual domain quote (€156.75 for 2 years, €6.53/month), which replaced an earlier €1-2/month estimate.

Chosen Pricing

Actual candidate pricing: **€3/month**, or **€25/year** (effective **€2.08/month** — a ~31% discount for paying annually, a meaningfully bigger discount than the SaaS-standard "2 months free" convention, which is a reasonable trade at this subscriber scale: reduced churn and better cash flow predictability matter more than maximizing per-subscriber revenue when the base is small).

Billing	Effective €/month	Breakeven at €42–55/month (realistic)	Breakeven at €28–92/month (full range)
Monthly (€3)	€3.00	14–19 subscribers	10–31 subscribers
Annual (€25)	€2.08	21–27 subscribers	14–45 subscribers

A mix skewed toward annual subscribers needs a somewhat larger base to hit the same breakeven, in exchange for stickier, more predictable revenue — not a problem, just worth knowing going in. If both billing options are offered at checkout, the annual price should always display its effective monthly rate alongside the yearly total (e.g. "€25/year — €2.08/mo"), both for clarity and to stay on the right side of EU consumer-pricing transparency norms once real payments are involved.

Price/month	Breakeven at €42–55/month (realistic)	Breakeven at €28–92/month (full range)
€2	21–28 subscribers	14–46 subscribers
€3	14–19 subscribers	10–31 subscribers
€4	11–14 subscribers	7–23 subscribers
€5	9–11 subscribers	6–19 subscribers

Reading This

- At any of these four price points, breakeven sits in the single or low double digits of paying subscribers — a genuinely low bar, and one that's realistic to clear even with a small, engaged niche audience (e.g. from the Future of Work vertical) rather than needing mass-market reach.
- Higher price points reach breakeven faster (fewer subscribers needed), but the trade-off is conversion: a lower price is easier to get people to say yes to, so the right choice depends on which side of that trade-off matters more to you, not on the breakeven math alone.
- This calculation ignores payment processing fees (Stripe and similar typically take roughly 1.5% + a small fixed fee per transaction) and any free tier that doesn't convert to paid — both push the real breakeven point slightly higher than shown, but not by enough to change the order of magnitude at these prices.

- It also doesn't account for your own time (maintenance, spot-checks) or recoup the one-time build cost from Cost of Building, above — this is breakeven on running costs only, not on the project as a whole.
- A useful way to read the range: if you land near the low end of running costs (€42/month) and price at €4–€5, you only need roughly 9–11 paying subscribers to cover hosting and API costs entirely — everything beyond that is margin.

Legal & Compliance Considerations

Not legal advice — these are the areas worth raising with a lawyer before launch, given the aggregator model and an EU (Malta) base. None of them are reasons not to build this; they're compliance work to fold into the build, most of it one-time setup rather than ongoing burden.

EU Press Publishers' Right (DSM Directive, Article 15)

The most specific legal risk to this exact business model. EU law gives press publishers a right over online use of their publications by news aggregators and media-monitoring services — the exact category this product falls into. The right does **not** cover mere hyperlinking or "individual words or very short extracts," which is why the Copyright Note and Story Deduplication sections above already call for fresh, non-copied summaries rather than reproduced snippets — that posture is the right one and substantially reduces this risk. But "very short extract" is a genuinely contested legal boundary (it triggered a real dispute between Google News and French publishers, resulting in a competition-law fine), so the summaries need to stay meaningfully rewritten and short, not just technically reworded versions of a source's own phrasing.

Scraping & Text-Data-Mining Rules

EU copyright law allows text-and-data-mining (which covers automated scraping/ingestion) for any purpose, including commercial, **unless the rightsholder has expressly reserved their rights via a machine-readable opt-out** (commonly a robots.txt directive or explicit terms-of-service language). The ingestion worker's scraper should check and respect robots.txt and each source's stated ToS before scraping it — preferring RSS feeds and official APIs (as the Source Map above already does for most sources) sidesteps this almost entirely, since those are offered for exactly this kind of reuse. Scraping a site that has explicitly opted out is a contract/copyright risk, not typically a criminal one, but worth avoiding rather than testing.

GDPR

Collecting email addresses for the digest sign-up means processing personal data under GDPR, which applies fully given the EU base and EU readers. Practically this means: a clear privacy policy, consent-based (not pre-ticked) sign-up, a data processing agreement with whichever email provider is used (Resend/Beehiiv — both are set up to offer these as standard), and honoring data subject rights (access, deletion) on request. This is standard, well-trodden compliance work for any small SaaS — not a special risk of this project, but a real setup checklist item.

AI Act Transparency (Article 50)

Since the site's summaries and classifications are AI-generated, the EU AI Act's transparency obligations (in force since 2 August 2026) point toward disclosing that AI-generated content is AI-generated — a simple, visible note (e.g. "summaries generated by AI, reviewed by spot-check") is the practical fix and is also good practice independent of the legal requirement, especially given your own professional focus on responsible AI governance.

VAT (EU One-Stop-Shop)

Worth flagging early because the threshold is lower than it looks: EU-established businesses owe VAT in the customer's country once **combined cross-border B2C digital-service sales reach €10,000/year** (below that, home-country Malta VAT rate applies). At €2–€5/month, roughly 170–420 paying subscribers would cross that threshold within a year — not an immediate concern at breakeven-level subscriber counts, but worth registering for the One-Stop-Shop scheme once growth approaches that range, since it's a single quarterly filing rather than per-country registration.

Editorial / Reputational Risk

The tone axis categorizes the **story** — which bucket of AI news coverage this is (progress, practical value, failure, or harm) — not a verdict on the company or person the story happens to involve. Tagging a story "Ugly" is closer to how a "Crime" section categorizes a type of story than to an accusation against a named party; the categorization itself carries little reputational risk. The residual risk sits elsewhere: in the underlying facts being wrong or the fresh summary misrepresenting what a source actually reported, which is a factual-accuracy problem rather than a tone-labeling one. The spot-check queue and the requirement that summaries stay faithful to source facts (Story Deduplication & Grouping, above) are the actual safeguards here, more so than the tone tag itself.

Trademark / Attribution

Naming and linking to sources (Reuters, FT, Bloomberg, etc.) as factual attribution is standard aggregator practice and not a trademark issue. Using their logos or visual branding without permission would be — stick to text links and source names.